Secure Multiparty Computation
Sprint 3

Developer | Hasnain Abdur Rehman| hasnain@bu.edu
Developer | Pierre-François Wolfe | pwolfe@bu.edu
Developer | Samyak Jain | samyakj@bu.edu
Developer | Suli Hu | sulihu@bu.edu
Developer | Yufeng Lin | yflin@bu.edu
Mentor/Client | John Liagouris | liagos@bu.edu
Mentor/Client | Vasiliki Kalavri | vkalavri@bu.edu
Subject-Matter Expert | Mayank Varia | varia@bu.edu

BOSTON UNIVERSITY

Presentation Outline

- Project Recap
- Project Goals & Sprint 3 Stories/Tasks
- Work Accomplished & Information Learned
  - Cloudlab/Bare-Metal Progress
  - Geni-Scripts & Bare-Metal results
  - Docker progress
- Project Organization Assessment (Burndown)
- Sprint 4 goals (Mentor priorities)

Recap of Multi-Party Computation (MPC)

- MPC enables...
  - Shared Computation on Private Data
  - Protects the Privacy of Data
  - Mutually Agreed Computation
- Our mentors…
  - Are using three party Secret Sharing MPC
  - Perform Relational Queries with MPC
  - Keep all parts secure vs. splitting into secure and insecure steps
- Our mission…
  - Profile this new MPC library
  - Identify bottlenecks
  - Compare deployment scenarios and find the best performance

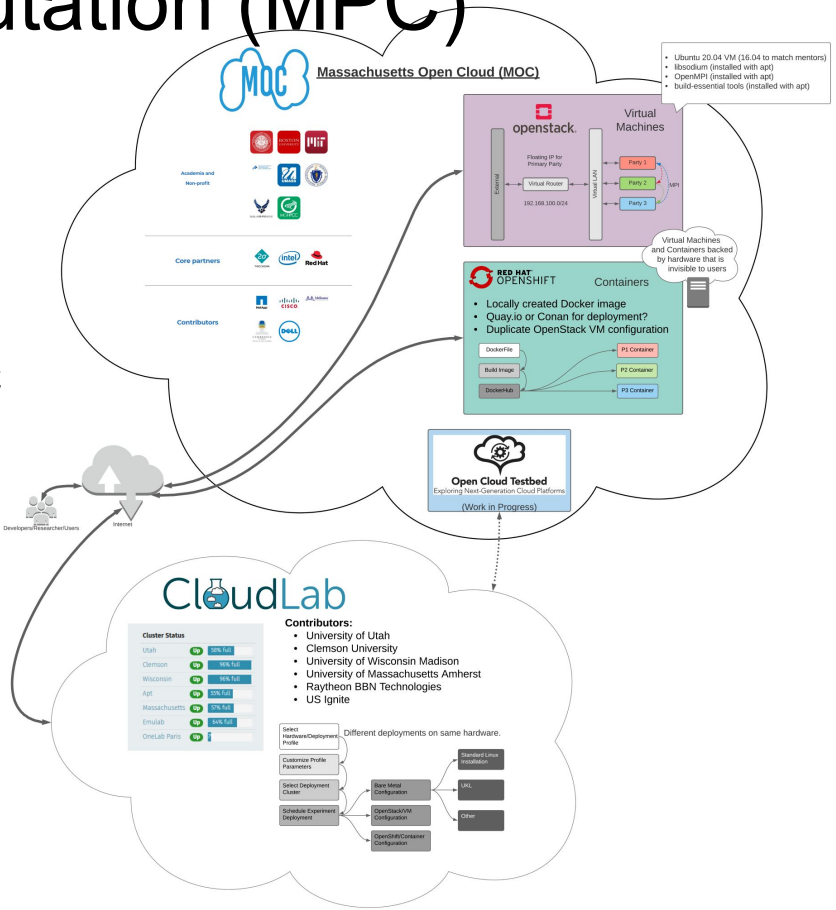

Project Goals & Sprint 3 Stories/Tasks

- VMs -
  - Improve existing instrumentation (delayed)
  - Explore tracing and profiling outputs from Score-p (delayed)
- Containers -
  - MPC test on single-container deployment (runs but debugging)
  - MPC test on multi-container deployment (waiting for single-container debug)
- Bare Metal -
  - MPC test on multi-bare-metal deployment (runs)
  - Initial data for exp-exchange collected, plotted
  - Challenges: Need to test on different profile deployment on bare metal
- Other -
  - Preparation for paper presentation (progress made, final revisions remain)

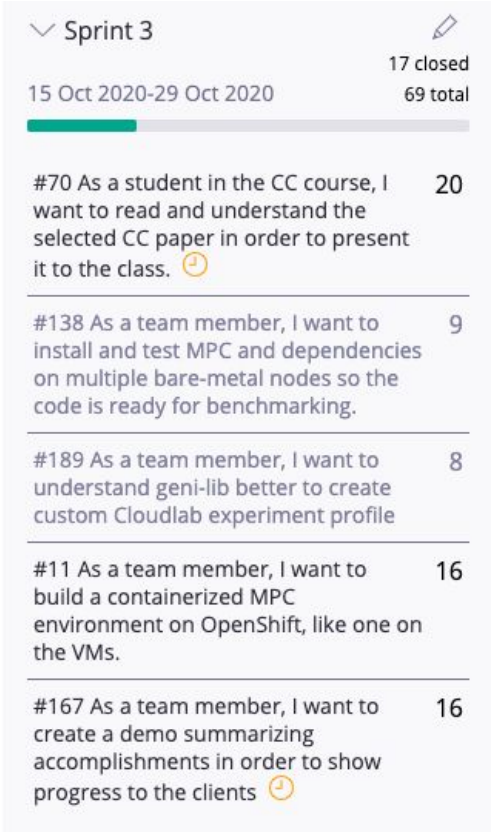

BOSTON UNIVERSITY

Change of Plans

- Challenges
  - Time spent on paper presentation
  - MOC downtime
  - CloudLab reservation challenges
  - Docker debugging
  - Team time conflicts this sprint (travel, exams, ...)
- Adjustment
  - Focus on running exp-exchange
    - Bare-metal on CloudLab → geni scripts
    - Containers → Local tests → Docker debugging

CloudLab/Bare-Metal Progress

- Specify new testing profile/environment
- Install dependencies needed for multi-nodes MPI
- Identify different testing environment
- Resolve communication between multi-nodes
- Test MPI functionality
- Run exp-exchange

ID	Node	Type	Status	Startup	Image	SSH command (if you provided your own key)	Actions
node-0	c220g1-030823	c220g1	ready	n/a	emulab-ops/UBUNTU14-64-STD	ssh -p 22 yflin@c220g1-030823.wisc.cloudlab.us	
node-1	c220g1-031103	c220g1	ready	n/a	emulab-ops/UBUNTU14-64-STD	ssh -p 22 yflin@c220g1-031103.wisc.cloudlab.us	
node-2	c220g1-031102	c220g1	ready	n/a	emulab-ops/UBUNTU14-64-STD	ssh -p 22 yflin@c220g1-031102.wisc.cloudlab.us	
client	c220g1-030826	c220g1	ready	n/a	emulab-ops/UBUNTU14-64-STD	ssh -p 22 yflin@c220g1-030826.wisc.cloudlab.us	
switch	c220g1-031113	c220g1	ready	n/a	emulab-ops/UBUNTU14-64-STD	ssh -p 22 yflin@c220g1-031113.wisc.cloudlab.us	

BOSTON UNIVERSITY

Custom CloudLab Experiments

- CloudLab Options
  - “Jacks” GUI
  - Hand crafted GENI RSpec (XML)
  - geni-lib → RSpec
- geni-lib script
  - Generates RSpec files with Python script
  - Much more readable
- Some Parameters
  - Number, type of nodes
  - Size of nodes
  - Type of link between nodes
  - Physical hardware

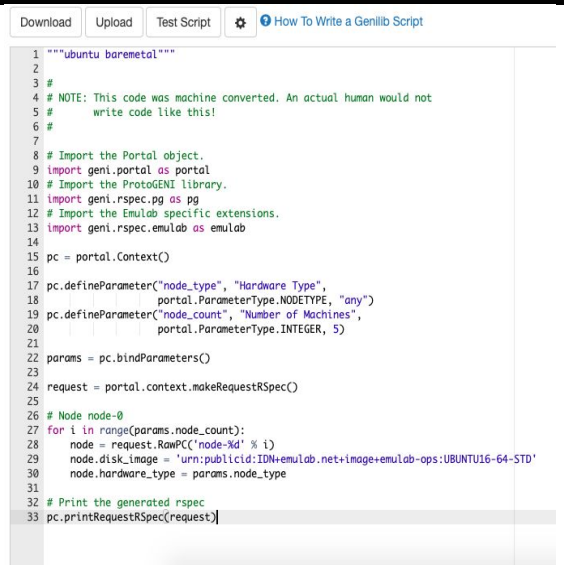

Custom Cloudlab Profile (geni-lib script)

Generated RSpec

Some Bare-Metal Results

- Cloudlab resources weren't available
  - Started new experiment on ARM based m400 nodes (only available nodes)
  - Used 'centos-n-bare-metal' profile as reference
  - Changed geni script to run with Ubuntu 16.04 OS
- Exp-exchange test
  - MPI batched, MPI eager (sync, async)
- Bare Metal tests faster than VMs
  - Due to ARM (on bare metal) vs x86 (on VMs)??
- Notes:
  - Sent huge traffic over shared control network on Cloudlab
  - Will create own LAN system to avoid this

MPC running on Bare Metal

```
SJain@node-0:~$ cd ccproject/experiments/
SJain@node-0:~/ccproject/experiments$ mpirun --host localhost,128.110.152.72,128
.110.152.47 -np 3 exp-exchange 1000
--------------------------------------------------------------------------
[[5240,1],1]: A high-performance Open MPI point-to-point messaging module
was unable to find any relevant network interfaces:

Module: OpenFabrics (openib)
  Host: node-1

Another transport will be used instead, although this may result in
lower performance.
--------------------------------------------------------------------------
BATCHED 1000    0.00004
SYNC    1000    0.02564
ASYNC   1000    0.02559
[node-0.mpc.mpcproject-pg0.utah.cloudlab.us:15666] 2 more processes have sent he
lp message help-mpi-btl-base.txt / btl:no-nics
[node-0.mpc.mpcproject-pg0.utah.cloudlab.us:15666] Set MCA parameter "orte_base_
help_aggregate" to 0 to see all help / error messages
SJain@node-0:~/ccproject/experiments$
```

Mentor MOC VM vs Our VM vs Bare Metal

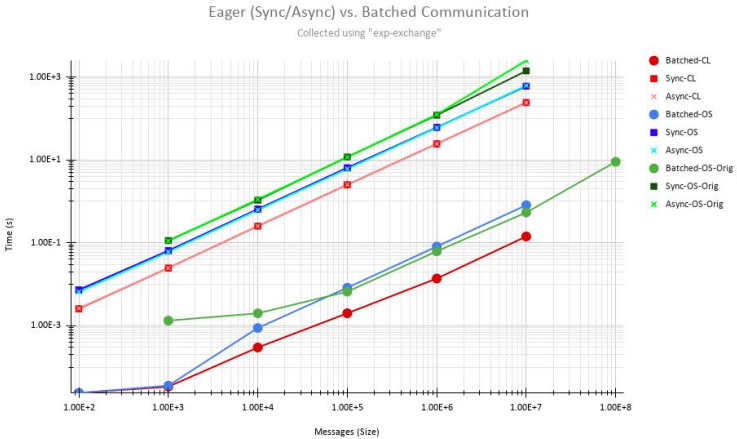

Docker as a Stepping Stone to OpenShift

- Plan for Testing
  - Local single-container MPC test
  - Local multi-container MPC test
  - Some type of more automated grouping
    - Local Kubernetes?
  - Deploy to OpenShift
- Execution of Plan
  - Discovered Interesting Docker Features
  - Worked around some OpenMPI issues
  - Still debugging/evaluating some messages from OpenMPI

Local Docker Setup

- Previously on Windows 10
  - Non-Pro → Docker Toolbox (deprecated)
  - Pro → Docker Desktop (using Hyper-V)
- Now on Windows 10
  - Anyone → Docker Desktop (WSL2)

1. Features

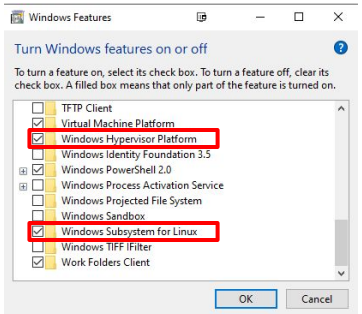

2. Linux

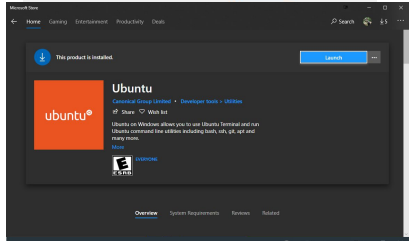

3. Docker

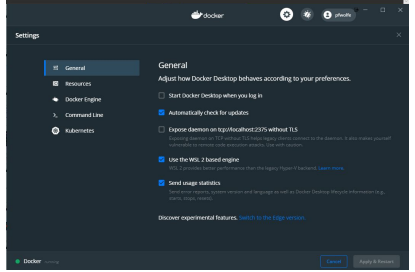

4. Terminal

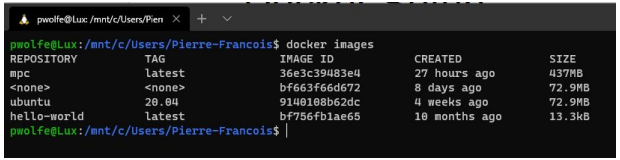

Docker in WSL2

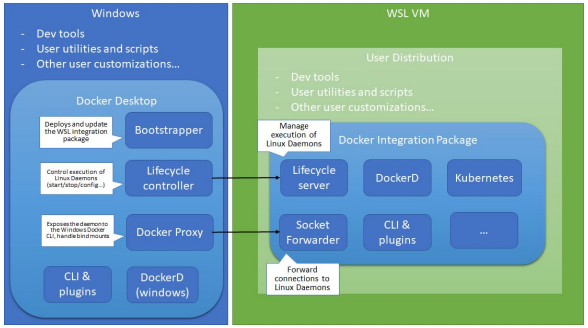

WSL Distros & Versions

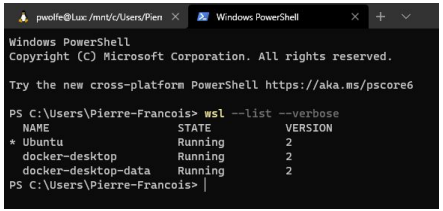

Image/reference: https://code.visualstudio.com/blogs/2020/03/02/docker-in-wsl2

Dockerfile with MPC Dependencies

- Interactive testing
- Dockerfile creation
  - Reference MPI Dockerfile
  - Insights from interactive Docker
- OpenMPI
  - Prevents running as root
  - Solution 1:
    - --allow-run-as-root
  - Solution 2: (per: https://github.com/open-mpi/ompi/pull/5597)
    - OMPI_ALLOW_RUN_AS_ROOT=1
    - OMPI_ALLOW_RUN_AS_ROOT_CONFIRM=1
- Future Steps:
  - We already minimize RUN commands
  - Multi-stage build to copy only needed binaries: see https://docs.docker.com/develop/develop-images/multistage-build/

```
# Dockerfile based on: https://github.com/oweidner/docker.openmpi/blob/master/Dockerfile
# Build this image: docker build -t mpc .

FROM ubuntu:20.04

MAINTAINER Pierre-Francois Wolfe <pwolfe@bu.edu>

ENV USER mpc

ENV HOME=/home/${USER}

ARG DEBIAN_FRONTEND=noninteractive

RUN apt update -y && \
    apt-get install -y --no-install-recommends sudo apt-utils && \
    apt-get install -y --no-install-recommends openssh-server \
    make gcc libopenmpi-dev openmpi-bin libsodium23 libsodium-dev && \
    apt clean && \
    apt purge && \
    rm -rf /var/lib/apt/lists/* /tmp/* /var/tmp/*

RUN useradd -ms /bin/bash mpc

COPY src/* /home/${USER}/src/
COPY experiments/* /home/${USER}/experiments/
COPY tests/* /home/${USER}/tests/

WORKDIR /home/${USER}/experiments
RUN make exp-exchange

ENV OMPI_ALLOW_RUN_AS_ROOT=1
ENV OMPI_ALLOW_RUN_AS_ROOT_CONFIRM=1

CMD mpirun -np 3 exp-exchange 1000
```


BOSTON UNIVERSITY

MPC in Docker Container Debugging...

- Running tests.sh → OK
- Running exp-exchange
  - Issue with size greater than 505… ex: with 1000
  - Cryptic message…
- Determine source…
  - Some clues but overall meaning still unclear

```
mpirun -np 3 exp-exchange1000
# which produces the following:
root@ebd1c7f24dfe:~/experiments#
[ebd1c7f24dfe:00046] Read -1, expected 8000, errno = 1
[ebd1c7f24dfe:00046] Read -1, expected 8000, errno = 1
[ebd1c7f24dfe:00047] Read -1, expected 8000, errno = 1
[ebd1c7f24dfe:00047] Read -1, expected 8000, errno = 1
[ebd1c7f24dfe:00045] Read -1, expected 8000, errno = 1
[ebd1c7f24dfe:00047] Read -1, expected 8000, errno = 1
[ebd1c7f24dfe:00046] Read -1, expected 8000, errno = 1
BATCHED 1000    0.00005
SYNC    1000    0.00047
ASYNC   1000    0.00042
```

Size changes value

[Hostname:PID]

BOSTON UNIVERSITY

Debugging Technique

Rank 0 → main party
Rank 1,2 → parties 2,3

- Code snippet entry point
  - Modify rank to attach to specific process
  - pid will print
- Attach with GDB using pid
- Change val i to continue
- Step through program
- Identify functions:
  - generate_and_share_random_data
  - exchange_rsz_seeds
  - exchange_shares_array
- MPI_Send/MPI_Recv pairs
  - 4, 1, and 2 respectively

```
# From: https://www.open-mpi.org/faq/?category=debugging
If (rank == 0) {
    volatile int i = 0;
    char hostname[256];
    gethostname(hostname, sizeof(hostname));
    printf("PID %d on %s ready for attach\n", getpid(), hostname);
    fflush(stdout);
    while (0 == i)
        sleep(5);
}
```

Changing rank to inspect different parties → message appears when the MPI_Recv is evaluated..

BOSTON UNIVERSITY

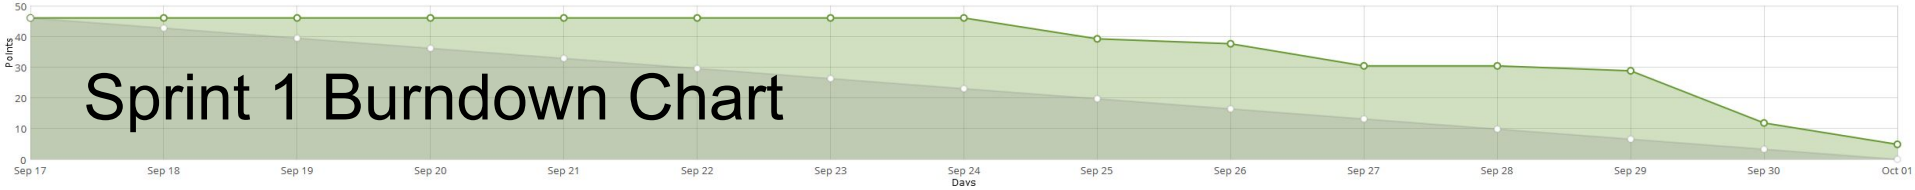

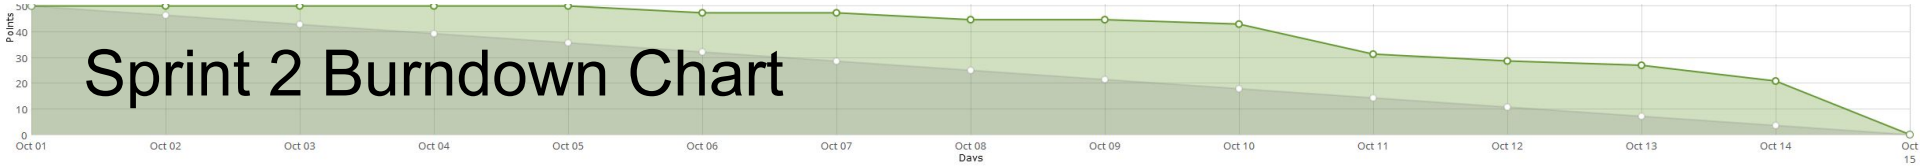

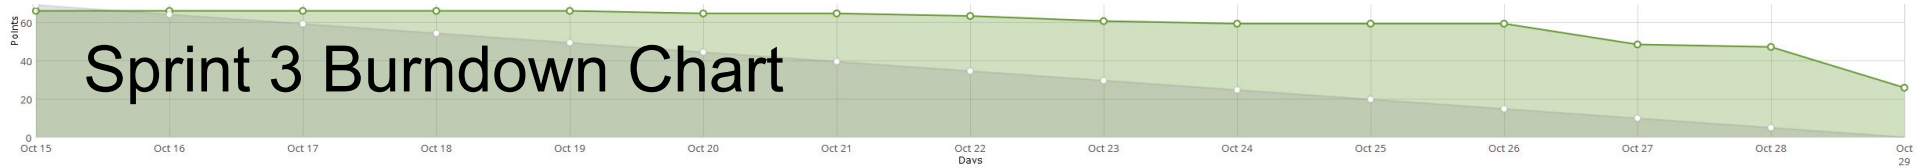

BOSTON UNIVERSITY

Sprint 4 - Some Known Stories

- VMs -
  - As a researcher, I want to improve the existing test instrumentation in order to more easily collect extra data samples, especially for large message sizes.
  - As a team member, I want to further explore the tracing and profiling outputs from Score-p to determine how to best assess performance.
- Containers -
  - As a team member, I want to fix issues with OpenMPI when running exp-exchange in docker to move closer to a local multi-container test and OpenShift deployment.
- Bare Metal -
  - As a team member, I want to make some improvements to my geni-lib script in order to refine my test environment and be able to capture more exp-exchange data runs
  - As a team member, I want to employ my geni-lib insights to create custom environments for testing OpenStack and OpenShift on CloudLab

Thank you

...any questions?

BOSTON UNIVERSITY

Backup Slides

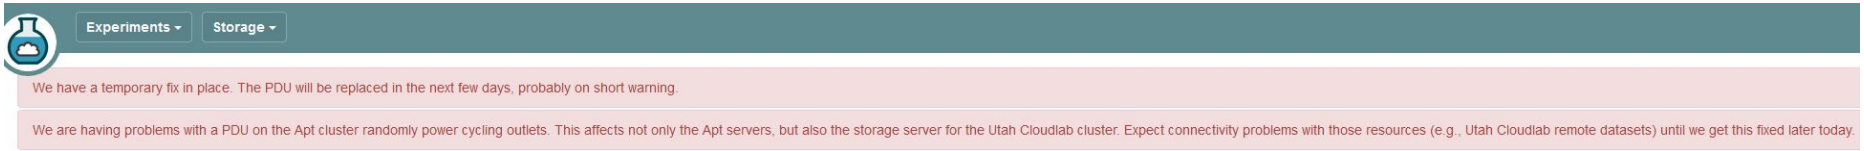

It seems there is an ongoing issue with CloudLab